

# Car Price Prediction Regression

## Data Loading and Preprocessing:

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.simplefilter("ignore")
```

```
In [2]: cp = pd.read_csv("C:\PGA32\MeriSkill\Oasis Infobyte\car data.csv")
```

```
In [3]: cp
```

Out[3]:

	Car_Name	Year	Selling_Price	Present_Price	Driven_kms	Fuel_Type	Selling_type	Transmission	Owner
<b>0</b>	ritz	2014	3.35	5.59	27000	Petrol	Dealer	Manual	0
<b>1</b>	sx4	2013	4.75	9.54	43000	Diesel	Dealer	Manual	0
<b>2</b>	ciaz	2017	7.25	9.85	6900	Petrol	Dealer	Manual	0
<b>3</b>	wagon r	2011	2.85	4.15	5200	Petrol	Dealer	Manual	0
<b>4</b>	swift	2014	4.60	6.87	42450	Diesel	Dealer	Manual	0
...	...	...	...	...	...	...	...	...	...
<b>296</b>	city	2016	9.50	11.60	33988	Diesel	Dealer	Manual	0
<b>297</b>	brio	2015	4.00	5.90	60000	Petrol	Dealer	Manual	0
<b>298</b>	city	2009	3.35	11.00	87934	Petrol	Dealer	Manual	0
<b>299</b>	city	2017	11.50	12.50	9000	Diesel	Dealer	Manual	0
<b>300</b>	brio	2016	5.30	5.90	5464	Petrol	Dealer	Manual	0

301 rows × 9 columns

In [4]: `cp.shape`

Out[4]: (301, 9)

In [5]: `cp.size`

Out[5]: 2709

In [6]: `cp.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Car_Name        301 non-null    object
1   Year            301 non-null    int64
2   Selling_Price    301 non-null    float64
3   Present_Price    301 non-null    float64
4   Driven_kms       301 non-null    int64
5   Fuel_Type        301 non-null    object
6   Selling_type     301 non-null    object
7   Transmission     301 non-null    object
8   Owner           301 non-null    int64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

In [7]: `cp.dtypes`

```
Out[7]: Car_Name      object
Year          int64
Selling_Price  float64
Present_Price  float64
Driven_kms     int64
Fuel_Type      object
Selling_type   object
Transmission   object
Owner          int64
dtype: object
```

In [8]: `cp.describe().T`

```
Out[8]:
```

	count	mean	std	min	25%	50%	75%	max
<b>Year</b>	301.0	2013.627907	2.891554	2003.00	2012.0	2014.0	2016.0	2018.0
<b>Selling_Price</b>	301.0	4.661296	5.082812	0.10	0.9	3.6	6.0	35.0
<b>Present_Price</b>	301.0	7.628472	8.642584	0.32	1.2	6.4	9.9	92.6
<b>Driven_kms</b>	301.0	36947.205980	38886.883882	500.00	15000.0	32000.0	48767.0	500000.0
<b>Owner</b>	301.0	0.043189	0.247915	0.00	0.0	0.0	0.0	3.0

In [9]: `cp.corr()`

Out[9]:

	Year	Selling_Price	Present_Price	Driven_kms	Owner
<b>Year</b>	1.000000	0.236141	-0.047192	-0.524342	-0.182104
<b>Selling_Price</b>	0.236141	1.000000	0.878914	0.029187	-0.088344
<b>Present_Price</b>	-0.047192	0.878914	1.000000	0.203618	0.008058
<b>Driven_kms</b>	-0.524342	0.029187	0.203618	1.000000	0.089216
<b>Owner</b>	-0.182104	-0.088344	0.008058	0.089216	1.000000

## No Missing Value

In [10]: `cp.isnull().sum()`

Out[10]:

Car_Name	0
Year	0
Selling_Price	0
Present_Price	0
Driven_kms	0
Fuel_Type	0
Selling_type	0
Transmission	0
Owner	0

dtype: int64

In [11]: `set(cp.duplicated())`

Out[11]: {False, True}

In [12]: `cp.duplicated().sum()`

Out[12]: 2

```
In [13]: # Check for duplicates across all columns
duplicated = cp.duplicated()

# Print the number of duplicated instances
print("Number of duplicated instances:", duplicated.sum())

# Print the duplicated instances
cp[duplicated]
```

Number of duplicated instances: 2

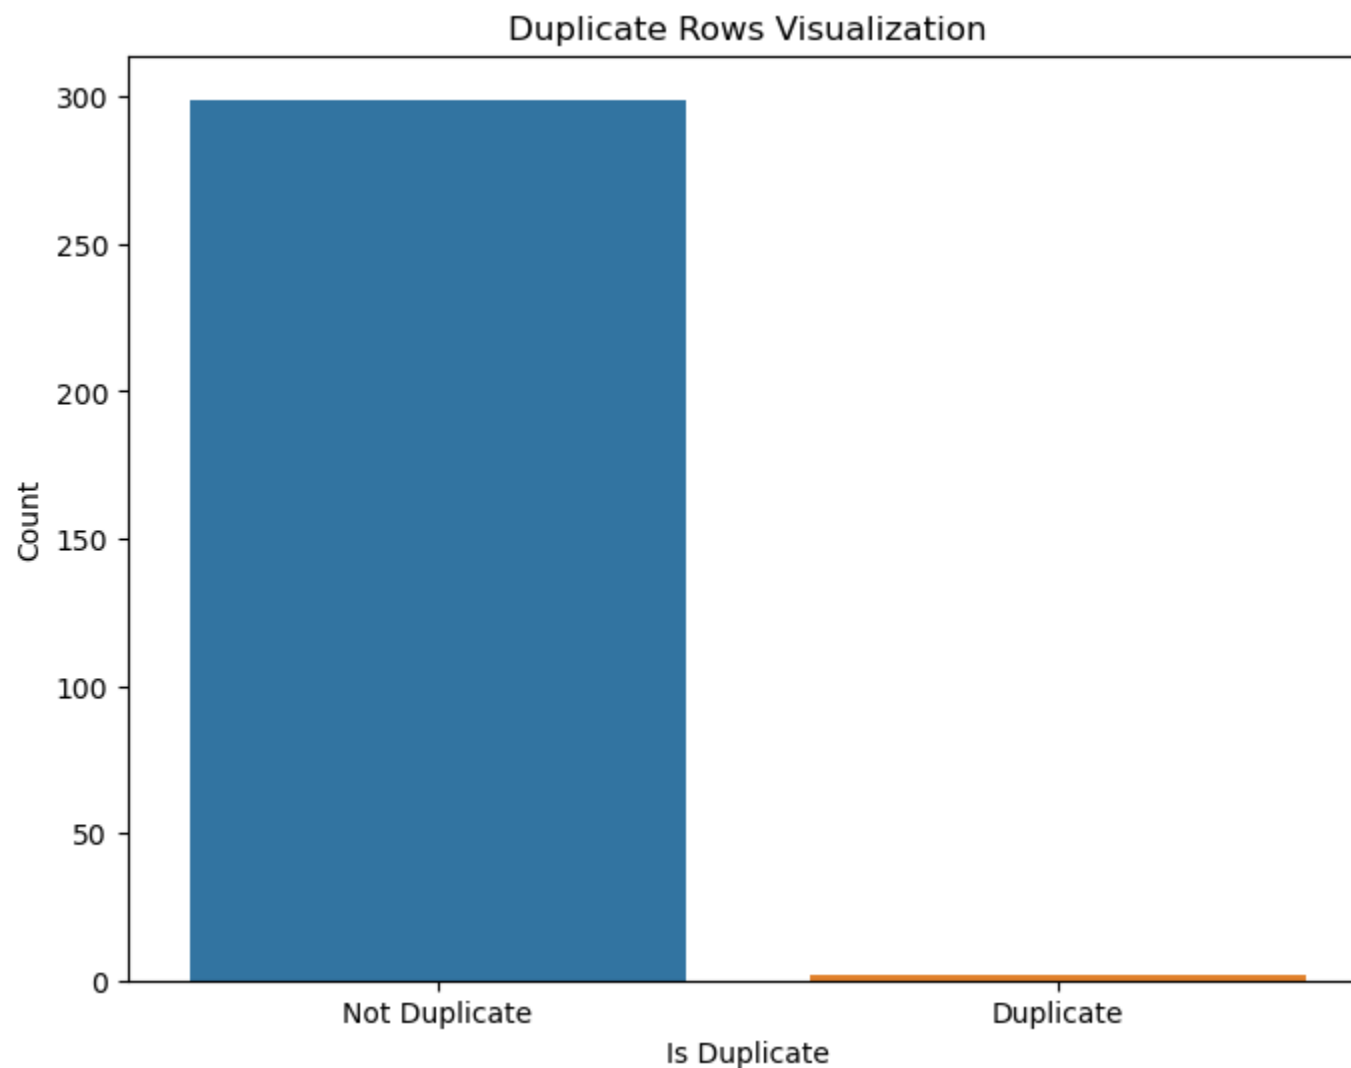
Out[13]:

	Car_Name	Year	Selling_Price	Present_Price	Driven_kms	Fuel_Type	Selling_type	Transmission	Owner
17	ertiga	2016	7.75	10.79	43000	Diesel	Dealer	Manual	0
93	fortuner	2015	23.00	30.61	40000	Diesel	Dealer	Automatic	0

```
In [14]: cp['Is_Duplicate'] = cp.duplicated()

# Create a count plot to visualize duplicates
plt.figure(figsize=(8, 6))
sns.countplot(x='Is_Duplicate', data=cp)
plt.xlabel('Is Duplicate')
plt.ylabel('Count')
plt.title('Duplicate Rows Visualization')
plt.xticks([0, 1], ['Not Duplicate', 'Duplicate'])
plt.show()

# Drop the "Is_Duplicate" column if not needed
cp.drop(columns=['Is_Duplicate'], inplace=True)
```



## Duplicate Dropping

```
In [15]: cp = cp.drop_duplicates(subset=['Car_Name', 'Year', 'Selling_Price', 'Present_Price', 'Driven_kms', 'Fuel_Type',  
                                         'Selling_type', 'Transmission', 'Owner'])
```

```
In [16]: # Check for duplicates across all columns  
         duplicated = cp.duplicated()  
  
         # Print the number of duplicated instances
```

```
print("Number of duplicated instances:", duplicated.sum())

# Print the duplicated instances
cp[duplicated]
```

Number of duplicated instances: 0

Out[16]:

Car_Name	Year	Selling_Price	Present_Price	Driven_kms	Fuel_Type	Selling_type	Transmission	Owner
----------	------	---------------	---------------	------------	-----------	--------------	--------------	-------

## Numerical Columns

In [17]: `cp.dtypes[cp.dtypes!="object"]`

Out[17]:

```
Year          int64
Selling_Price float64
Present_Price float64
Driven_kms    int64
Owner         int64
dtype: object
```

In [18]:


```
# Select numerical columns
numerical_columns = ['Year', 'Selling_Price', 'Present_Price', 'Driven_kms', 'Owner']

# Create a DataFrame containing only the numerical columns
numerical_df = cp[numerical_columns]
```

In [19]:

```
# Calculate the correlation matrix for numerical columns
correlation_matrix = numerical_df.corr()

# Create a heatmap
plt.figure(figsize=(15, 7))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap (Numerical Columns)')
plt.show()
```




```
In [20]: correlation_matrix.corr()
```

```
Out[20]:
```

	Year	Selling_Price	Present_Price	Driven_kms	Owner
Year	1.000000	0.199905	-0.178871	-0.910902	-0.464685
Selling_Price	0.199905	1.000000	0.921894	-0.203017	-0.576554
Present_Price	-0.178871	0.921894	1.000000	0.128255	-0.388411
Driven_kms	-0.910902	-0.203017	0.128255	1.000000	0.135237
Owner	-0.464685	-0.576554	-0.388411	0.135237	1.000000

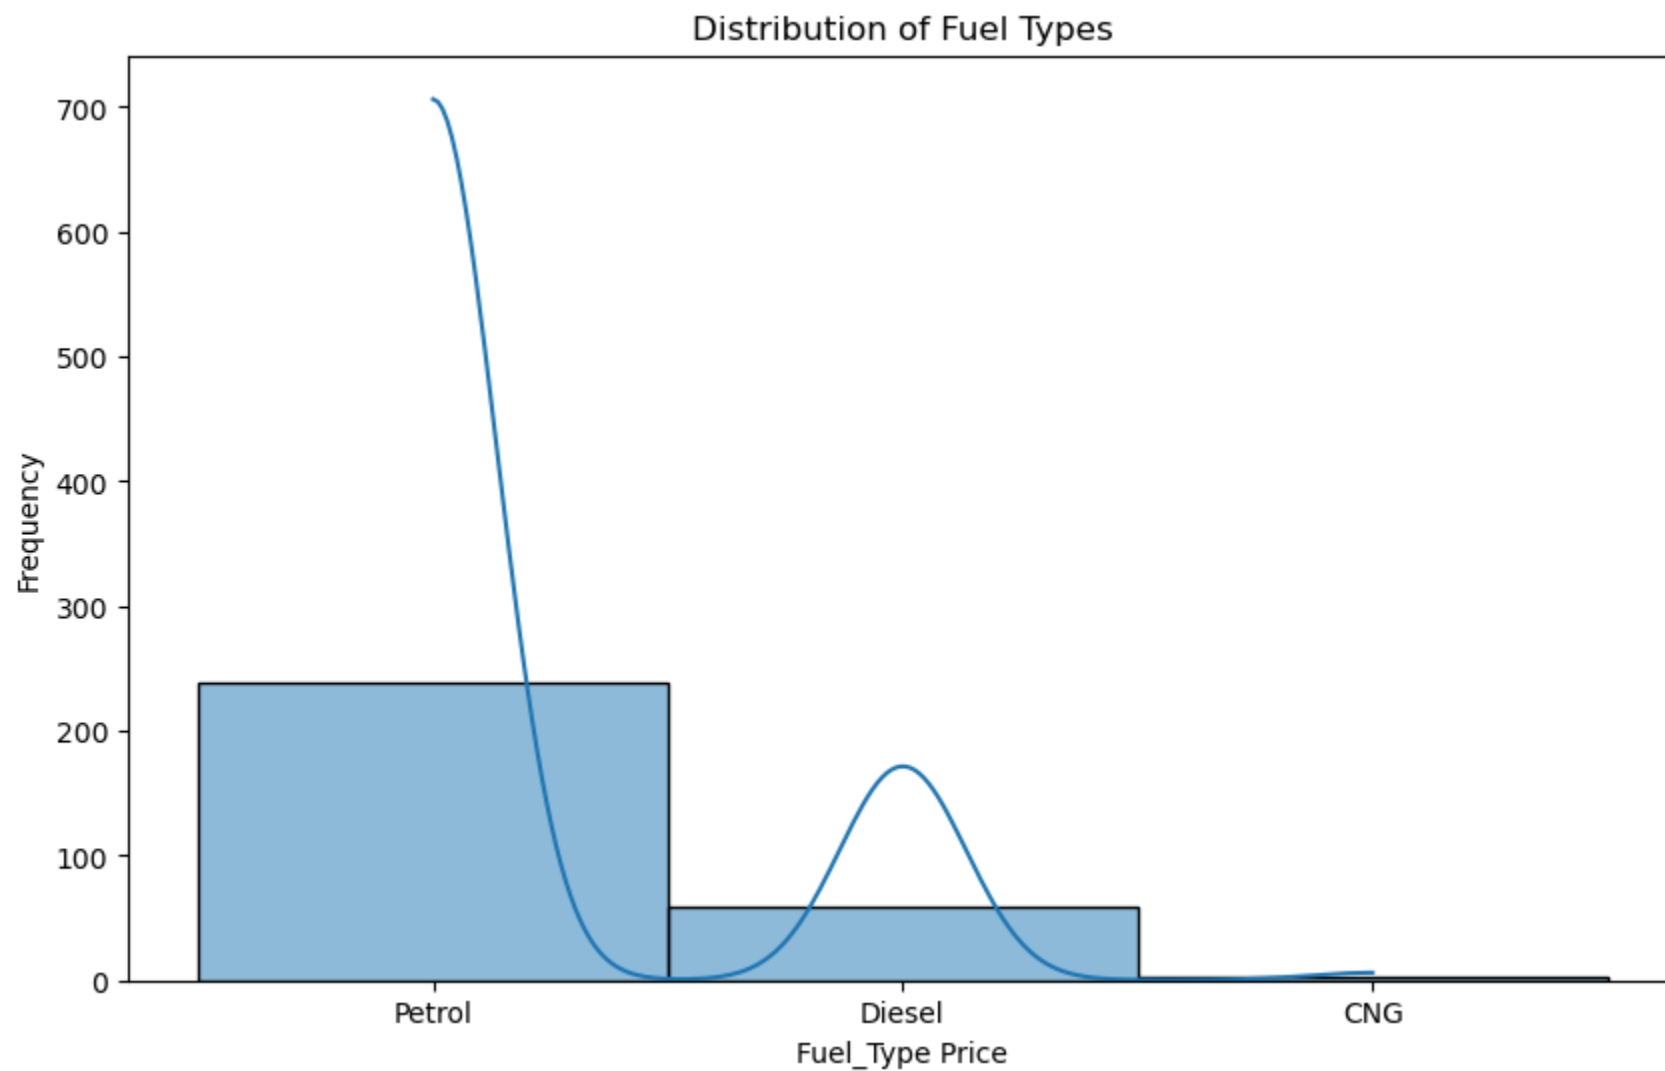


```
In [21]: plt.figure(figsize=(10, 6))
sns.histplot(cp['Selling_Price'], bins=20, kde=True)
plt.xlabel('Selling Price')
plt.ylabel('Frequency')
plt.title('Distribution of Selling Prices')
plt.show()
```




```
In [22]: plt.figure(figsize=(10, 6))
sns.histplot(cp['Fuel_Type'], bins=20, kde=True)
plt.xlabel('Fuel_Type Price')
```


```
plt.ylabel('Frequency')  
plt.title('Distribution of Fuel Types')  
plt.show()
```



```
In [23]: plt.figure(figsize=(10, 6))  
sns.histplot(cp['Selling_type'], bins=20, kde=True)  
plt.xlabel('Selling types')  
plt.ylabel('Frequency')  
plt.title('Distribution of Selling Types')  
plt.show()
```




```
In [24]: species_counts = cp['Fuel_Type'].value_counts()
plt.figure(figsize=(8, 8))
plt.title('Distribution of Fuel Types')
plt.pie(species_counts, labels=species_counts.index, autopct='%1.1f%%', startangle=140)
plt.axis('equal')
plt.show()
```




```
In [25]: species_counts = cp['Transmission'].value_counts()  
plt.figure(figsize=(8, 8))  
plt.title('Distribution of Transmissions Types')
```


```
plt.pie(species_counts, labels=species_counts.index, autopct='%1.1f%%', startangle=140)  
plt.axis('equal')  
plt.show()
```



```
In [26]: def univariate_numerical(data,var,graph_plot=True):  
    missing=data[var].isnull().sum()  
    min_n=data[var].min()  
    max_n=data[var].max()  
    var_n=data[var].var()  
    std_n=data[var].std()  
    p10=data[var].quantile(.1)  
    p25=data[var].quantile(.25)  
    p50=data[var].quantile(.5)  
    p75=data[var].quantile(.75)  
    p99=data[var].quantile(.99)  
    iqr=p75-p25  
  
    if graph_plot==True:  
        sns.histplot(data[var])  
        plt.show()  
        sns.boxplot(y=data[var])  
        plt.show()  
  
    results={"missing":missing,"min":min_n,"max":max_n,"var":var_n,"std":std_n,  
            "p10":p10,"p25":p25,"p50":p50,"p75":p75,"p99":p99,}  
    return results
```

```
In [27]: univariate_numerical(data=cp, var="Selling_Price")
```







```
Out[27]: {'missing': 0,  
          'min': 0.1,  
          'max': 35.0,  
          'var': 24.842645166214016,  
          'std': 4.984239677845962,  
          'p10': 0.4,  
          'p25': 0.8500000000000001,  
          'p50': 3.51,  
          'p75': 6.0,  
          'p99': 23.009999999999999}
```

```
In [28]: univariate_numerical(data=cp, var="Year")
```









```
Out[28]: {'missing': 0,  
          'min': 2003,  
          'max': 2018,  
          'var': 8.391843056272588,  
          'std': 2.8968678009658273,  
          'p10': 2010.0,  
          'p25': 2012.0,  
          'p50': 2014.0,  
          'p75': 2016.0,  
          'p99': 2017.0}
```


```
In [29]: univariate_numerical(data=cp, var="Present_Price")
```






```
Out[29]: {'missing': 0,  
          'min': 0.32,  
          'max': 92.6,  
          'var': 73.38203894830627,  
          'std': 8.566331708981755,  
          'p10': 0.7060000000000001,  
          'p25': 1.2,  
          'p50': 6.1,  
          'p75': 9.84,  
          'p99': 35.96}
```


```
In [30]: univariate_numerical(data=cp, var="Driven_kms")
```






```
Out[30]: {'missing': 0,  
          'min': 500,  
          'max': 500000,  
          'var': 1522183517.6298168,  
          'std': 39015.17035243876,  
          'p10': 5940.0,  
          'p25': 15000.0,  
          'p50': 32000.0,  
          'p75': 48883.5,  
          'p99': 143103.5199999999}
```

```
In [31]: univariate_numerical(data=cp, var="Owner")
```





```
Out[31]: {'missing': 0,
          'min': 0,
          'max': 3,
          'var': 0.061861686606361384,
          'std': 0.24872009690887745,
          'p10': 0.0,
          'p25': 0.0,
          'p50': 0.0,
          'p75': 0.0,
          'p99': 1.0}
```

## Driven KMS Vs. Selling Price and Present Price


```
In [32]: plt.figure(figsize=(10, 6))

# Create line plots for 'alcohol' and 'residual sugar' against 'density'
sns.lineplot(x="Selling_Price", y="Driven_kms", data=cp, label='Selling Prices', marker='o')
sns.lineplot(x="Present_Price", y="Driven_kms", data=cp, label='Present Prices', marker='s')

plt.title("Sales vs. Selling Price and Present Price", fontsize=16)
```



```
plt.xlabel("Sales", fontsize=12)
plt.ylabel("Value", fontsize=12)
plt.legend()
plt.tight_layout()
plt.show()
```




```
In [33]: fig, axes = plt.subplots(nrows=4, ncols=3, figsize=(20, 20))

# Modify the color list to match the number of columns in your dataset
colors = [ '#491D8B', '#6929C4', '#8A3FFC', '#A56EFF',
            '#7D3AC1', '#AF4BCE', '#DB4CB2', '#EB548C',
            '#EC96E0', '#A2128E', '#E8D9F3', '#641811' ]
```

```
# Loop through each column in your wine_df dataset
for index, column in enumerate(cp.columns):
    if index < 12: # Limit the iteration to the number of subplots
        ax = axes.flatten()[index]
        ax.hist(cp[column], color=colors[index], label=column)
        ax.legend(loc="best")

plt.suptitle("Histograms", size=20)
plt.tight_layout()
plt.show()
```





There is a statistically significant relationship between the "Selling\_Price" and "Present\_Price" & "Selling\_Price" and "Year" variables

```
In [34]: from scipy import stats
figure=pd.crosstab(cp["Selling_Price"], cp["Present_Price"])
stats.chi2_contingency(figure)
```

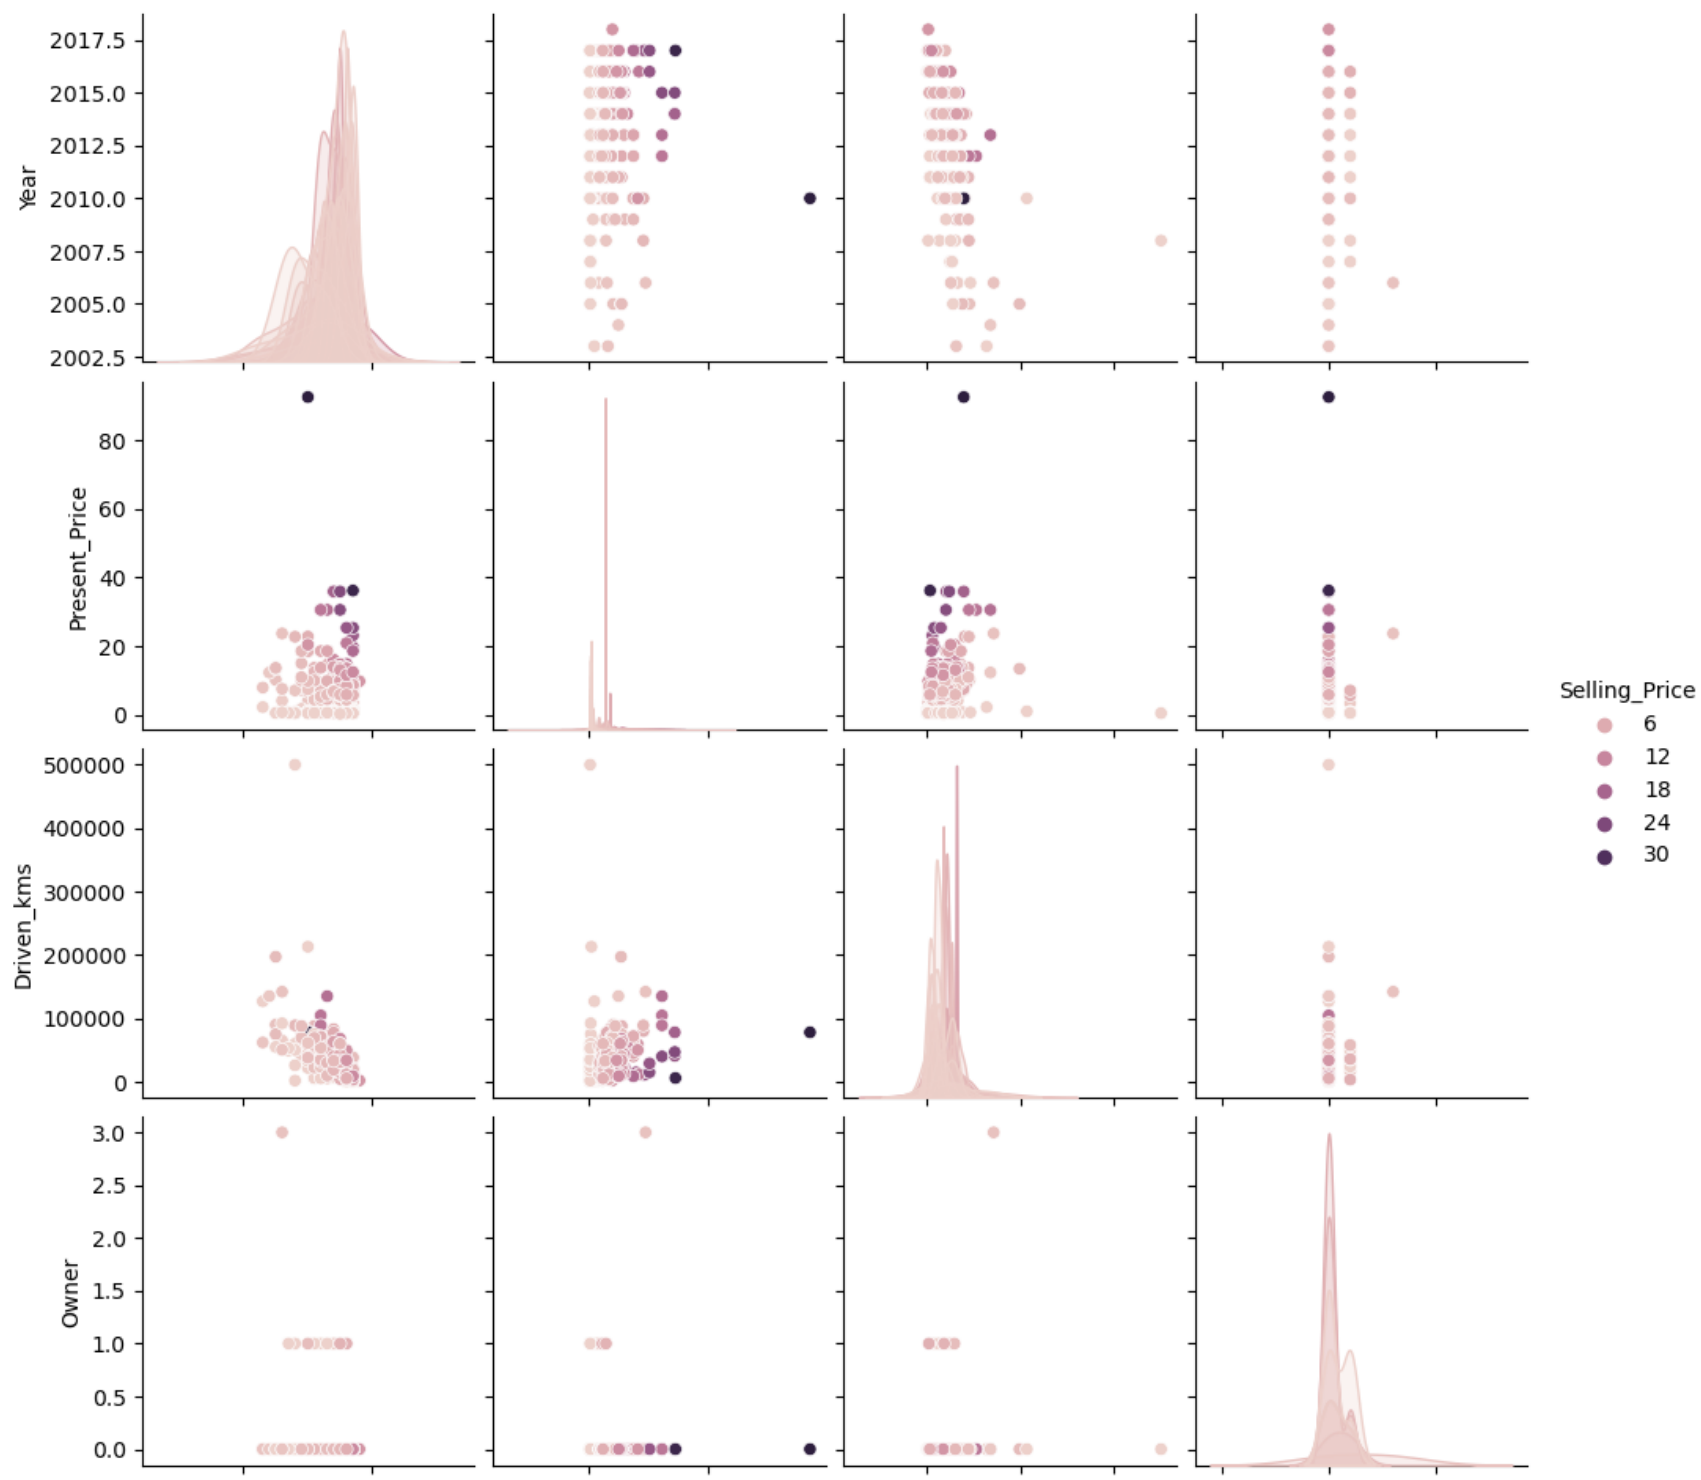
```
Out[34]: Chi2ContingencyResult(statistic=23320.2206292517, pvalue=0.006369652129717279, dof=22785, expected_freq=array([[0.00334448, 0.00334448, 0.00334448, ..., 0.01003344, 0.00334448, 0.00334448],
[0.00334448, 0.00334448, 0.00334448, ..., 0.01003344, 0.00334448, 0.00334448],
[0.00334448, 0.00334448, 0.00334448, ..., 0.01003344, 0.00334448, 0.00334448],
[0.00334448, 0.00334448, 0.00334448, ..., 0.01003344, 0.00334448, 0.00334448],
...,
[0.00334448, 0.00334448, 0.00334448, ..., 0.01003344, 0.00334448, 0.00334448],
[0.00334448, 0.00334448, 0.00334448, ..., 0.01003344, 0.00334448, 0.00334448],
[0.00334448, 0.00334448, 0.00334448, ..., 0.01003344, 0.00334448, 0.00334448],
[0.00334448, 0.00334448, 0.00334448, ..., 0.01003344, 0.00334448, 0.00334448]]))
```

```
In [35]: from scipy import stats
figure=pd.crosstab(cp["Selling_Price"], cp["Year"])
stats.chi2_contingency(figure)
```

```
Out[35]: Chi2ContingencyResult(statistic=2443.8008122656674, pvalue=0.04243557416005079, dof=2325, expected_freq=array([[0.00668896, 0.00334448, 0.01337793, ..., 0.1638796 , 0.11705686, 0.00334448], [0.00668896, 0.00334448, 0.01337793, ..., 0.1638796 , 0.11705686, 0.00334448], [0.00668896, 0.00334448, 0.01337793, ..., 0.1638796 , 0.11705686, 0.00334448], ..., [0.00668896, 0.00334448, 0.01337793, ..., 0.1638796 , 0.11705686, 0.00334448], [0.00668896, 0.00334448, 0.01337793, ..., 0.1638796 , 0.11705686, 0.00334448], [0.00668896, 0.00334448, 0.01337793, ..., 0.1638796 , 0.11705686, 0.00334448]]))
```

I am visualizing the correlation of the dataset with the seaborn library.

```
In [36]: sns.pairplot(cp, hue = "Selling_Price")  
plt.show()
```



## Correlation Matrix


```
In [37]: d=cp.corr()
def mylight(x):
    color="red" if abs(x)>=.6 else ""
    return "background-color: {}".format(color)
```

```
In [38]: d.style.applymap(mylight)
```

```
Out[38]:
```

	Year	Selling_Price	Present_Price	Driven_kms	Owner
Year	1.000000	0.234369	-0.053167	-0.525714	-0.181639
Selling_Price	0.234369	1.000000	0.876305	0.028566	-0.087880
Present_Price	-0.053167	0.876305	1.000000	0.205224	0.009948
Driven_kms	-0.525714	0.028566	0.205224	1.000000	0.089367
Owner	-0.181639	-0.087880	0.009948	0.089367	1.000000

```
In [39]: corrmatrix=cp.corr()
plt.figure(figsize=(15,8))
sns.heatmap(corrmatrix,cmap = 'plasma', annot=True)
plt.show()
```



There is a strong positive correlation (0.88) between "Selling\_Price" and "Present\_Price."

There is a strong negative correlation (-0.53) between "Driven\_kms" and "Year."

## Data Encoding:

```
In [40]: # checking the distribution of categorical data  
print(cp['Fuel_Type'].value_counts())
```



```
print(cp['Selling_type'].value_counts())  
print(cp['Transmission'].value_counts())
```

```
Petrol    239  
Diesel    58  
CNG        2  
Name: Fuel_Type, dtype: int64  
Dealer    193  
Individual 106  
Name: Selling_type, dtype: int64  
Manual    260  
Automatic  39  
Name: Transmission, dtype: int64
```

## Label Encoding

```
In [41]: # encoding "Fuel_Type" Column  
cp.replace({'Fuel_Type':{'Petrol':0,'Diesel':1,'CNG':2}},inplace=True)  
  
# encoding "Seller_Type" Column  
cp.replace({'Selling_type':{'Dealer':0,'Individual':1}},inplace=True)  
  
# encoding "Transmission" Column  
cp.replace({'Transmission':{'Manual':0,'Automatic':1}},inplace=True)  
cp
```

Out[41]:

	Car_Name	Year	Selling_Price	Present_Price	Driven_kms	Fuel_Type	Selling_type	Transmission	Owner
<b>0</b>	ritz	2014	3.35	5.59	27000	0	0	0	0
<b>1</b>	sx4	2013	4.75	9.54	43000	1	0	0	0
<b>2</b>	ciaz	2017	7.25	9.85	6900	0	0	0	0
<b>3</b>	wagon r	2011	2.85	4.15	5200	0	0	0	0
<b>4</b>	swift	2014	4.60	6.87	42450	1	0	0	0
...	...	...	...	...	...	...	...	...	...
<b>296</b>	city	2016	9.50	11.60	33988	1	0	0	0
<b>297</b>	brio	2015	4.00	5.90	60000	0	0	0	0
<b>298</b>	city	2009	3.35	11.00	87934	0	0	0	0
<b>299</b>	city	2017	11.50	12.50	9000	1	0	0	0
<b>300</b>	brio	2016	5.30	5.90	5464	0	0	0	0

299 rows × 9 columns

## Modeling: Feature Scaling

Pre-processing Steps for Machine Learning

```
In [42]: X = cp.drop(['Car_Name', 'Selling_Price'],axis=1)
          y = cp['Selling_Price']
```

```
In [43]: X
```

Out[43]:

	Year	Present_Price	Driven_kms	Fuel_Type	Selling_type	Transmission	Owner
<b>0</b>	2014	5.59	27000	0	0	0	0
<b>1</b>	2013	9.54	43000	1	0	0	0
<b>2</b>	2017	9.85	6900	0	0	0	0
<b>3</b>	2011	4.15	5200	0	0	0	0
<b>4</b>	2014	6.87	42450	1	0	0	0
...	...	...	...	...	...	...	...
<b>296</b>	2016	11.60	33988	1	0	0	0
<b>297</b>	2015	5.90	60000	0	0	0	0
<b>298</b>	2009	11.00	87934	0	0	0	0
<b>299</b>	2017	12.50	9000	1	0	0	0
<b>300</b>	2016	5.90	5464	0	0	0	0

299 rows × 7 columns

In [44]:

y

Out[44]:

```

0      3.35
1      4.75
2      7.25
3      2.85
4      4.60

```

```

...
296    9.50
297    4.00
298    3.35
299   11.50
300    5.30

```

Name: Selling\_Price, Length: 299, dtype: float64

## Splitting the data

```
In [45]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=13)
```

## Feature Scaling

```
In [46]: from sklearn.preprocessing import StandardScaler
```

```
In [47]: # Standardize the features
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

# Regression Models

## Linear Regression

```
In [48]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
```

```
In [49]: # Create and fit the linear regression model
lin_reg = LinearRegression()
```

```
In [50]: lin_reg.fit(X_train, y_train)
```

```
Out[50]: ▼ LinearRegression
LinearRegression()
```

```
In [51]: # Make predictions
y_pred = lin_reg.predict(X_test)
```

```
In [52]: # Evaluate the model
mse_linear = mean_squared_error(y_test, y_pred)
print("Linear Regression Mean Squared Error:", mse_linear)
```

Linear Regression Mean Squared Error: 2.405203022617024

```
In [53]: from sklearn.metrics import mean_squared_error, r2_score
```

```
In [54]: r2 = r2_score(y_test, y_pred)
print("R-squared (R^2) Score:", r2)
```

R-squared (R^2) Score: 0.8915921588461179

```
In [55]: # Optionally, you can print the model coefficients and intercept
print("Coefficients:", lin_reg.coef_)
```

```
print("Intercept:", lin_reg.intercept_)
```

Coefficients: [ 1.15565057 3.85889097 -0.23759324 0.53812924 -0.58201869 0.51898959  
-0.20249736]

Intercept: 4.456820083682

```
In [56]: #Check the test score and train score to the RandomForestRegressor algorithm
print(f'The Test_accuracy: {lin_reg.score(X_test,y_test)*100:.2f}')
```

```
#Train score for the data
```

```
print(f'The Train_accuracy: {lin_reg.score(X_train,y_train)*100:.2f}')
```

The Test\_accuracy: 89.16

The Train\_accuracy: 87.12

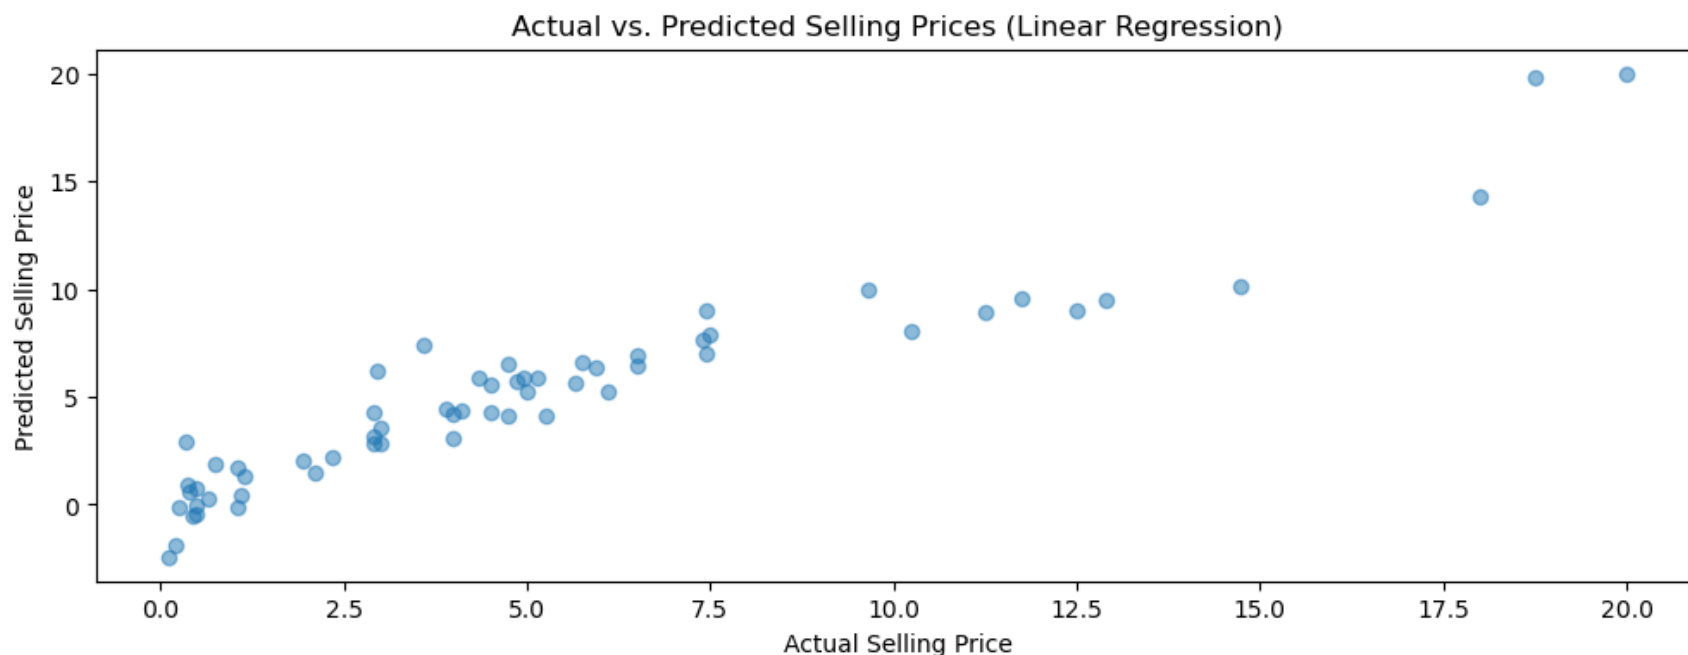
```
In [57]: print(f"Linear Regression Mean Squared Error (MSE): {mse_linear:.2f}")
print(f"Linear Regression R-squared (R2): {r2:.2f}")
```

Linear Regression Mean Squared Error (MSE): 2.41

Linear Regression R-squared (R2): 0.89

## Evaluation and Visualization:

```
In [58]: plt.figure(figsize=(12,4))
plt.scatter(y_test, y_pred, alpha=0.5)
plt.xlabel('Actual Selling Price')
plt.ylabel('Predicted Selling Price')
plt.title('Actual vs. Predicted Selling Prices (Linear Regression)')
plt.show()
```



## Decision Tree Regression

```
In [59]: from sklearn.tree import DecisionTreeRegressor
```

```
In [60]: Deci_Tree = DecisionTreeRegressor(random_state=17)
```

```
In [61]: # Train the model on the training set
Deci_Tree.fit(X_train, y_train)
```

```
Out[61]: ▼      DecisionTreeRegressor
DecisionTreeRegressor(random_state=17)
```

```
In [62]: # Make predictions on the test set
y_preds = Deci_Tree.predict(X_test)
```

```
In [63]: #Check the test score and train score to the RandomForestRegressor algorithm
print(f'Red Test_accuracy: {Deci_Tree.score(X_test,y_test)*100:.2f}')
```

```
#Train score for the data  
print(f'The Train_accuracy: {Deci_Tree.score(X_train,y_train)*100:.2f}')
```

The Test\_accuracy: 92.27  
The Train\_accuracy: 100.00

```
In [64]: # Evaluate the model  
mse = mean_squared_error(y_test, y_preds)  
r2 = r2_score(y_test, y_preds)
```


```
In [65]: print(f'Decision Tree Mean Squared Error: {mse:.2f}')
```

```
print(f'Decision Tree R-squared Score: {r2:.2f}')
```

Decision Tree Mean Squared Error: 1.72  
Decision Tree R-squared Score: 0.92

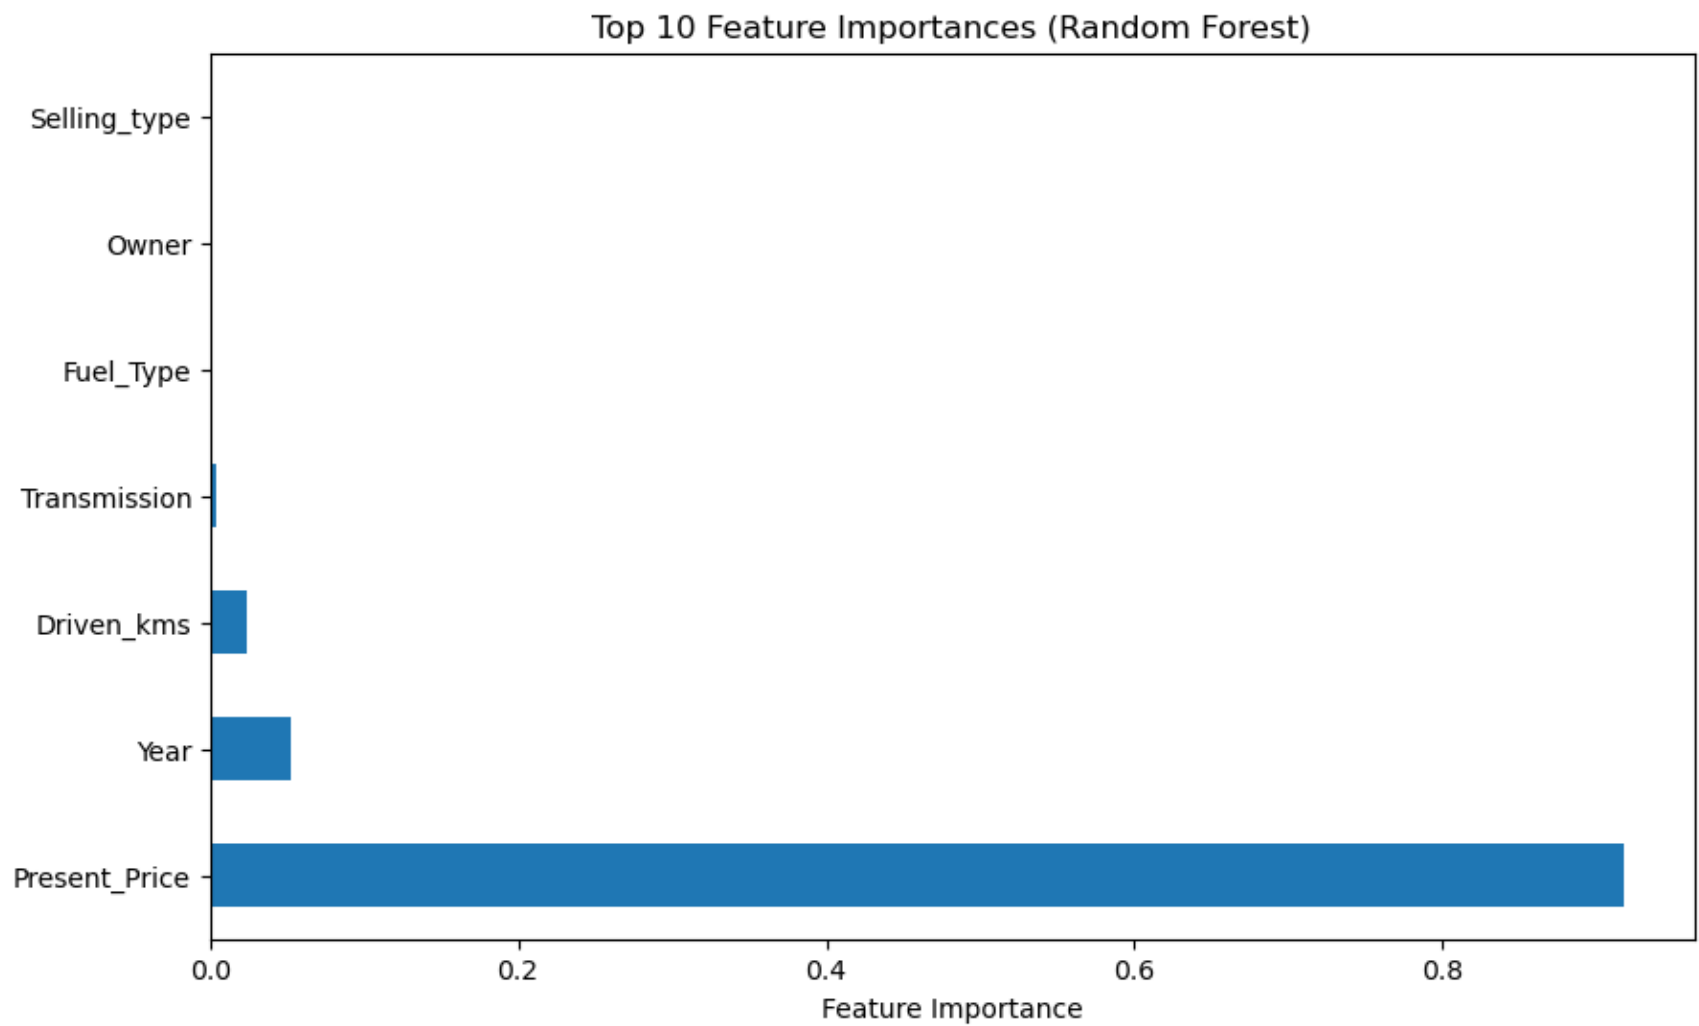
## Evaluation and Visualization:

```
In [66]: plt.figure(figsize=(12,4))  
  
plt.scatter(y_test, y_preds)  
plt.xlabel("Actual Price")  
plt.ylabel("Predicted Price")  
plt.title(" Actual Prices vs Predicted Prices(Desicision Tree Regression)")  
plt.show()
```



```
In [67]: feature_importances = pd.Series(Deci_Tree.feature_importances_, index=X.columns)
plt.figure(figsize=(10, 6))
feature_importances.nlargest(10).plot(kind='barh')
plt.xlabel('Feature Importance')
plt.title('Top 10 Feature Importances (Random Forest)')
plt.show()
```





## Support Vector Machine Regression Model

```
In [68]: from sklearn.svm import SVR
```

```
In [69]: # Initialize the Support Vector Regression model  
Support_V = SVR(kernel='linear', C=5)
```

```
In [70]: # Train the model on the scaled training set  
Support_V.fit(X_train, y_train)
```

Out[70]:

▼ SVR

SVR(C=5, kernel='linear')

```
In [71]: # Make predictions on the scaled test set
y_pred_s = Support_V.predict(X_test)
```

```
In [72]: #Check the test score and train score to the RandomForestRegressor algorithm
print(f'The Test_accuracy: {Support_V.score(X_test,y_test)*100:.2f}')
```

*#Train score for the data*

```
print(f'The Train_accuracy: {Support_V.score(X_train,y_train)*100:.2f}')
```

The Test\_accuracy: 87.81  
The Train\_accuracy: 84.10

```
In [73]: # Evaluate the model
mse = mean_squared_error(y_test, y_pred_s)
r2 = r2_score(y_test, y_pred_s)
```

```
In [74]: print(f'Support Vector Machine Mean Squared Error: {mse:.2f}')
```


```
print(f'Support Vector Machine R-squared Score: {r2:.2f}')
```

Support Vector Machine Mean Squared Error: 2.70  
Support Vector Machine R-squared Score: 0.88

## Evaluation and Visualization:

```
In [75]: plt.figure(figsize=(12,4))

plt.scatter(y_test, y_pred_s)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title(" Actual Prices vs Predicted Prices(Support Vector Machine Regression)")
plt.show()
```



## Random Forest Regression Model

```
In [76]: from sklearn.ensemble import RandomForestRegressor
```

```
In [77]: # Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [78]: # Initialize the Random Forest Regressor model
Random_Forest_model = RandomForestRegressor(n_estimators=100, random_state=17)
```

```
In [79]: # Train the model on the training set
Random_Forest_model.fit(X_train, y_train)
```

```
Out[79]: ▼      RandomForestRegressor
RandomForestRegressor(random_state=17)
```

```
In [80]: # Make predictions on the test set
y_pred_ss = Random_Forest_model.predict(X_test)
```

```
In [81]: #Check the test score and train score to the RandomForestRegressor algorithm
print(f'The Test_accuracy: {Random_Forest_model.score(X_test,y_test)*100:.2f}')

#Train score for the data
print(f'The Train_accuracy: {Random_Forest_model.score(X_train,y_train)*100:.2f}')
```

The Test\_accuracy: 96.39  
The Train\_accuracy: 98.27

```
In [82]: # Evaluate the model
mse = mean_squared_error(y_test, y_pred_ss)
r2 = r2_score(y_test, y_pred_ss)
```


```
In [83]: print(f'Random Forest Mean Squared Error: {mse:.2f}')
print(f'Random Forest R-squared Score: {r2:.2f}')
```

Random Forest Mean Squared Error: 0.77  
Random Forest R-squared Score: 0.96

## Evaluation and Visualization:

```
In [84]: plt.figure(figsize=(12,4))

plt.scatter(y_test, y_pred_ss)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title(" Actual Prices vs Predicted Prices(Random Forest Regression)")
plt.show()
```



The model with the highest test accuracy is the "Random Forest Regression Model" with a test accuracy of 96.39

The model with the highest R-squared score is the Random Forest with a score of 0.96. The Random Forest regression model is considered to have the best goodness of fit to the data.

```
In [87]: models = ["Linear Regression Model", "Decision Tree Regression Model", "SVM Regression Model", "Random Forest Regression Model"]
train_accuracies = [87.12, 100.00, 84.10, 98.27]
test_accuracies = [89.16, 92.27, 87.81, 96.39]
mse_accuracies = [2.41, 1.72, 2.70, 0.77]

# Set the width of the bars
bar_width = 0.25 # Adjust the bar width as needed

# Create an array of equally spaced positions for the bars
x = np.arange(len(models))


plt.figure(figsize=(10, 6))
```

```
# Create a grouped bar chart
plt.bar(x - bar_width, train_accuracies, bar_width, label='Train Accuracy')
plt.bar(x, test_accuracies, bar_width, label='Test Accuracy', alpha=0.8)
plt.bar(x + bar_width, mse_accuracies, bar_width, label='Mean Square Error Accuracy', alpha=0.8)

# Set title and labels
plt.title("Model Performance")
plt.xlabel("Models")
plt.ylabel("Accuracy")
plt.legend()

# Rotate x-axis labels for better readability
plt.xticks(x, models, rotation=45, ha="right")

# Display the plot
plt.tight_layout()
plt.show()
```



THANK YOU!